

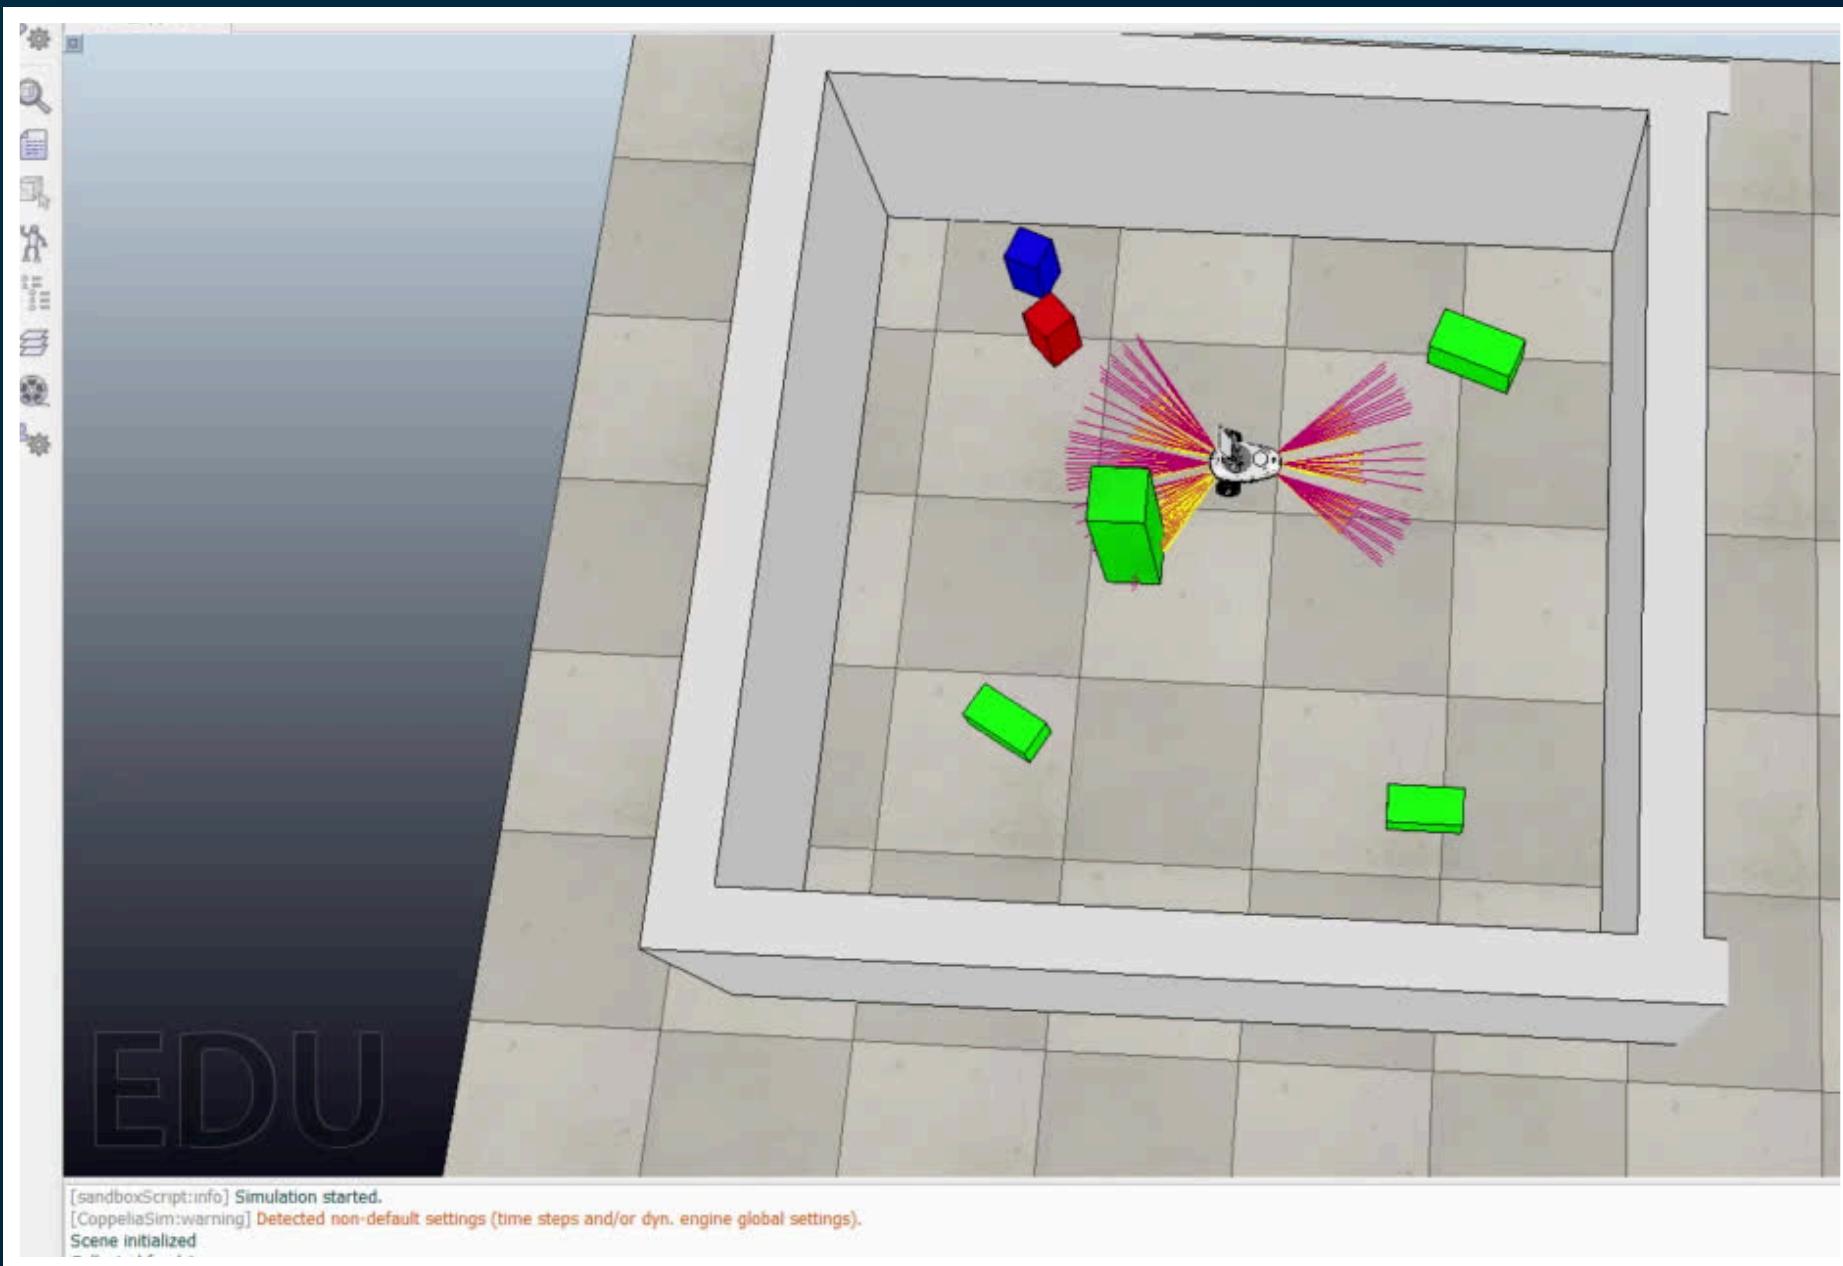


TASK 2 - FORAGING LEARNING MACHINES Robobo

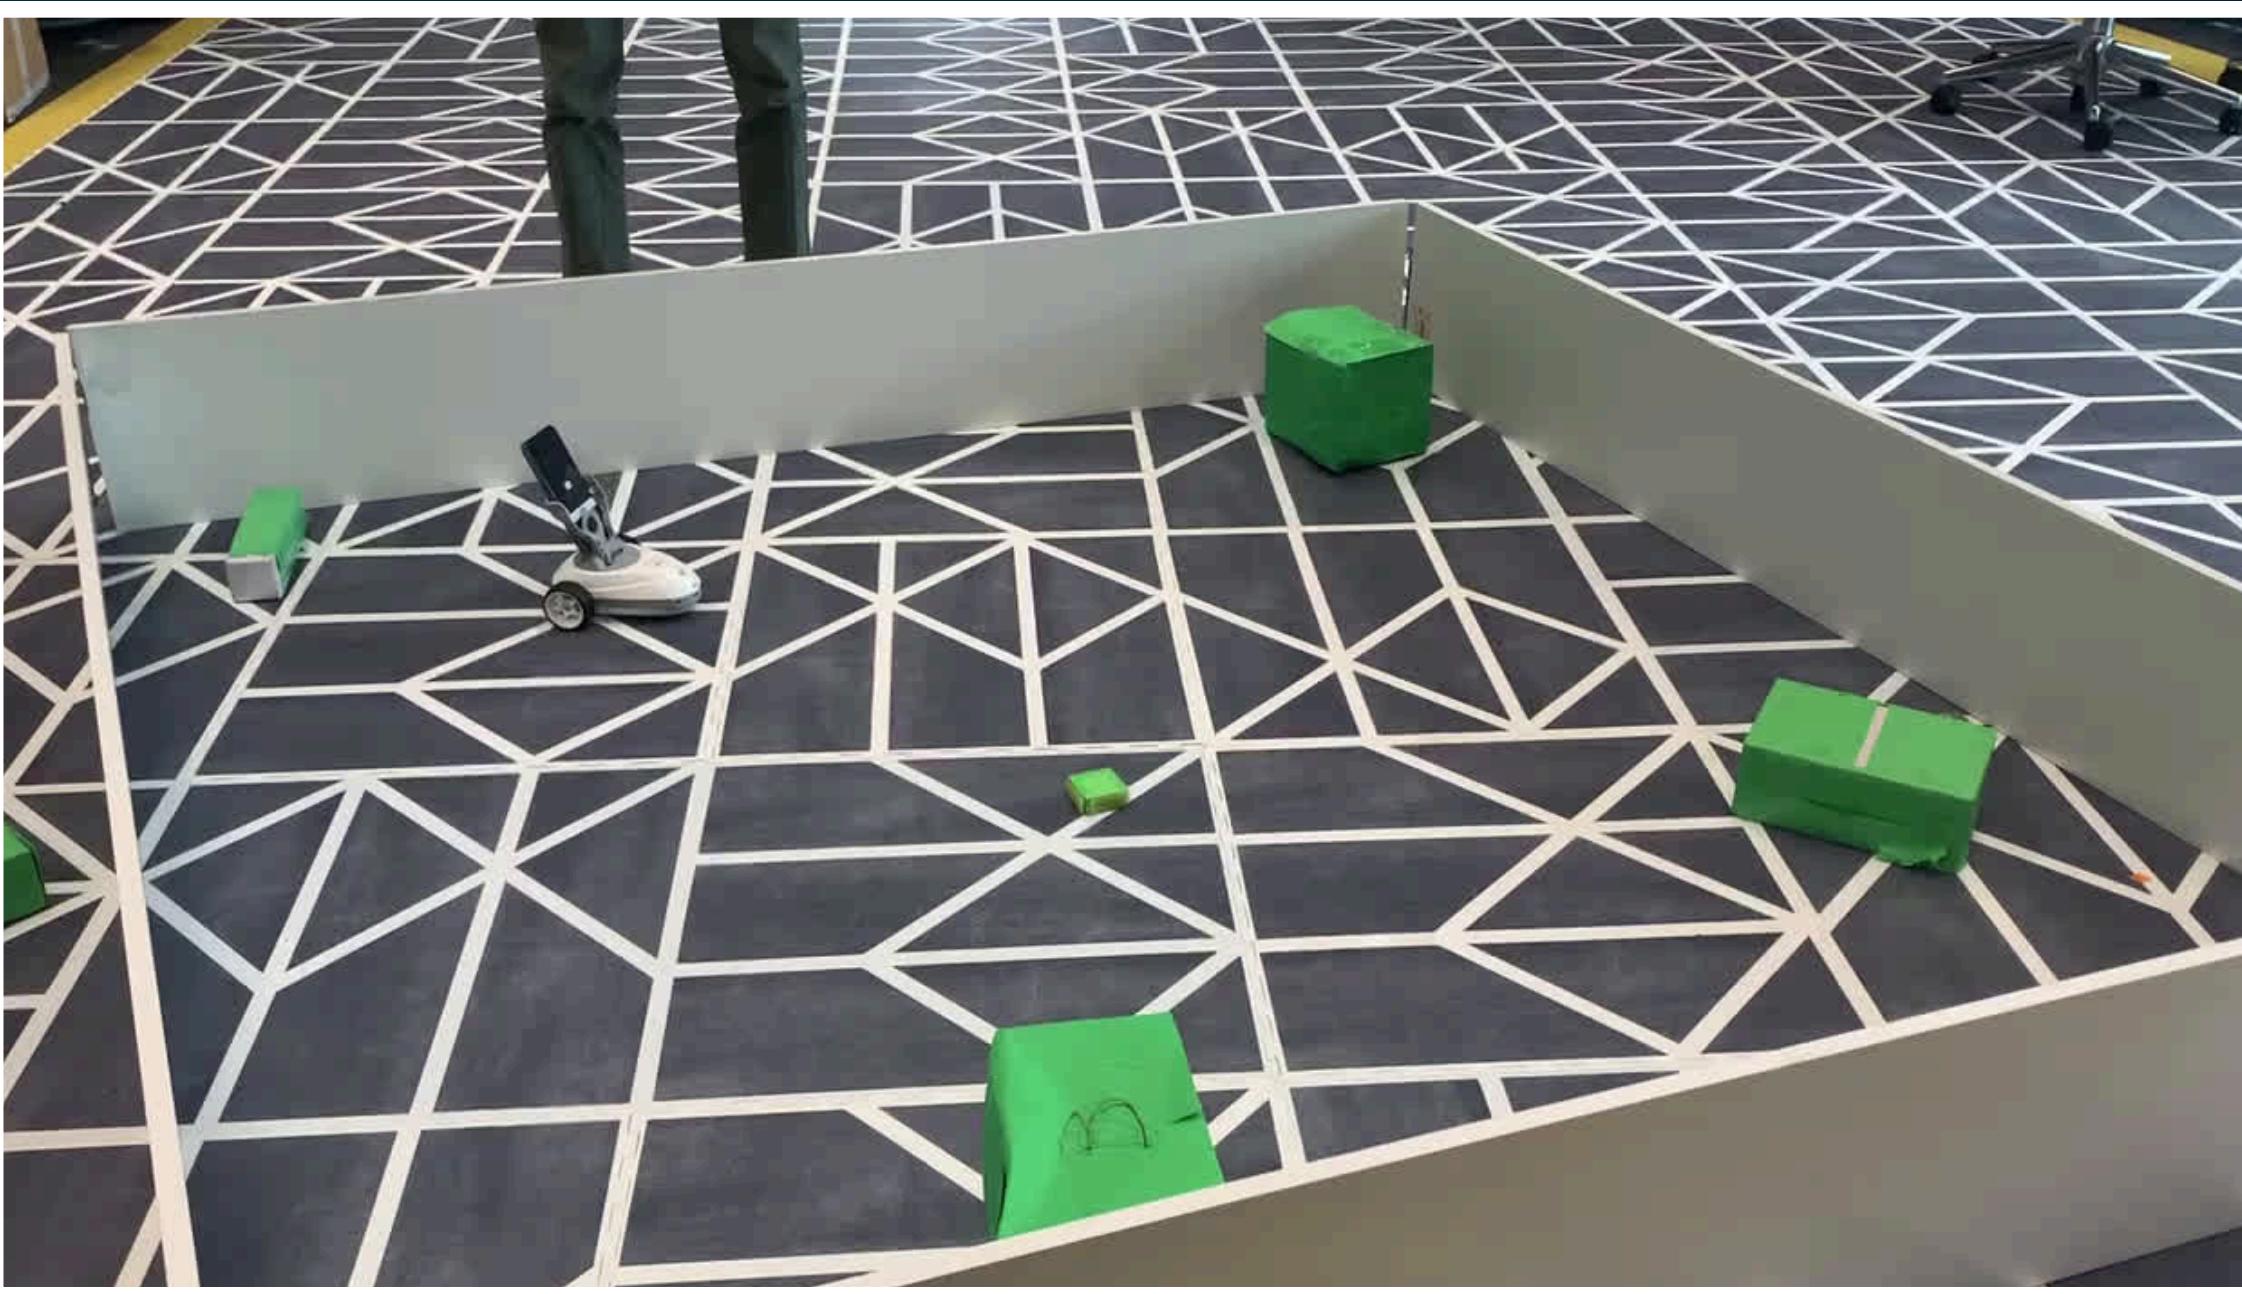
GROUP 14

Stergios Ntanavaras
Thijs Grootjans
Quirinus de Ruijter

SIMULATION DEMO



HARDWARE DEMO



PROBLEM & MOTIVATION

- **Objective:** Develop an AI-driven robot that learns to touch and collect the largest number of “food” items (packages) per minute within a given environment.
- **Challenges:**
 - Designing algorithms that can quickly and efficiently learn from interactions with the environment.
 - Requires advanced perception and adaptability.
 - Complex decision-making.
 - Solutions that work well in a virtual environment might not scale effectively to real-world scenarios.
 - Exogenous factors can significantly affect the performance of sensors like cameras.



PROBLEM & MOTIVATION

- **Utility and Importance to AI Community:**

- Pushes the development of algorithms, particularly in perception and real-time decision making.
- Leads to significant improvements in efficiency and cost reduction in various industries.
- Creates the foundations for autonomous systems operating in the real-world.
- Could revolutionize operations to industries like logistics, warehousing and manufacturing.



METHODOLOGY

- **Algorithm:** Q-learning, a model-free RL algorithm.
 - *Purpose:*
 - Develop a strategy for the robot to autonomously learn and optimize its actions in the environment.
 - *Goals:*
 - Enable the robot to learn from its own actions and experiences.
 - Improve the robot's actions to maximize the number of items touched and collected.
 - Refining decision-making over time.
 - *Learning Process:*
 - Initializes Q-values at zero.
 - Updates them based on rewards and future reward estimates:
 - $$Q(s, a) = Q(s, a) + \alpha [R(s, a) + \gamma \max Q'(s', a') - Q(s, a)]$$



METHODOLOGY

- **Object Detection:** Analyzing the distribution of green pixels by capturing and processing the robot's camera.
 - *Steps:*
 - Capture the current view from the robot's front-facing camera.
 - Clip the image to focus on the relevant part.
 - Calculate the percentage of green pixels : $(\#G/\#T) \times 100$
 - Analyze greenness distribution:
 - Dividing the image into 4 vertical sections and focusing on the middle 2 (center view).
 - Identifying the section with the highest number of green pixels.
 - *Goal:* Create a state that provides the robot a location of the nearest green object.



METHODOLOGY

- **Sensing and Reality Gap:**
 - Sensors used:
 - Front Central IR sensor & Mobile Camera.
 - Infrared Sensor Binning Logic:
 - Converting raw infrared sensor values into discrete bins based on the defined thresholds.
 - Greenness Binning Logic:
 - Categorizing the greenness percentage into discrete bins.
 - Different thresholds for simulation and hardware, due to Reality Gap.

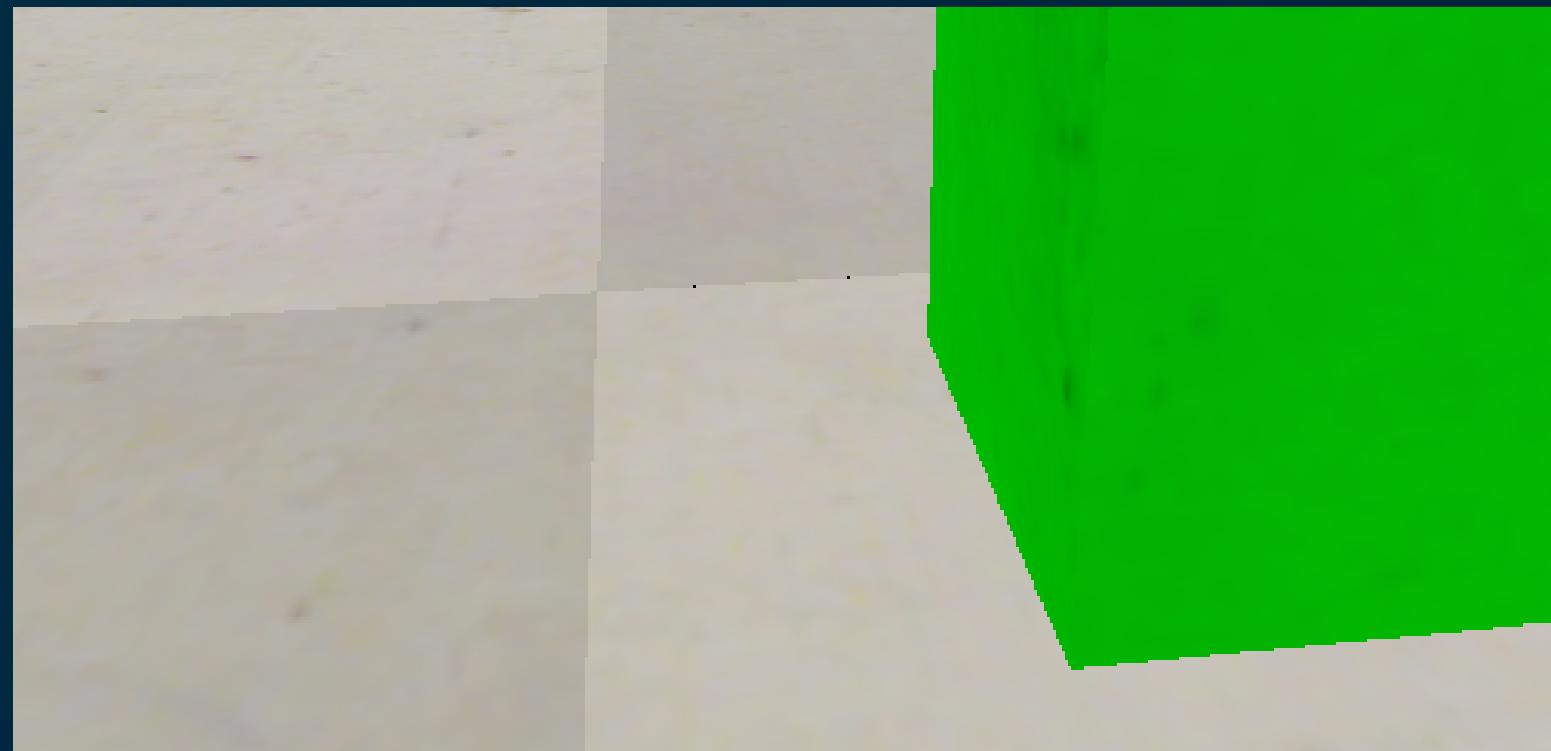
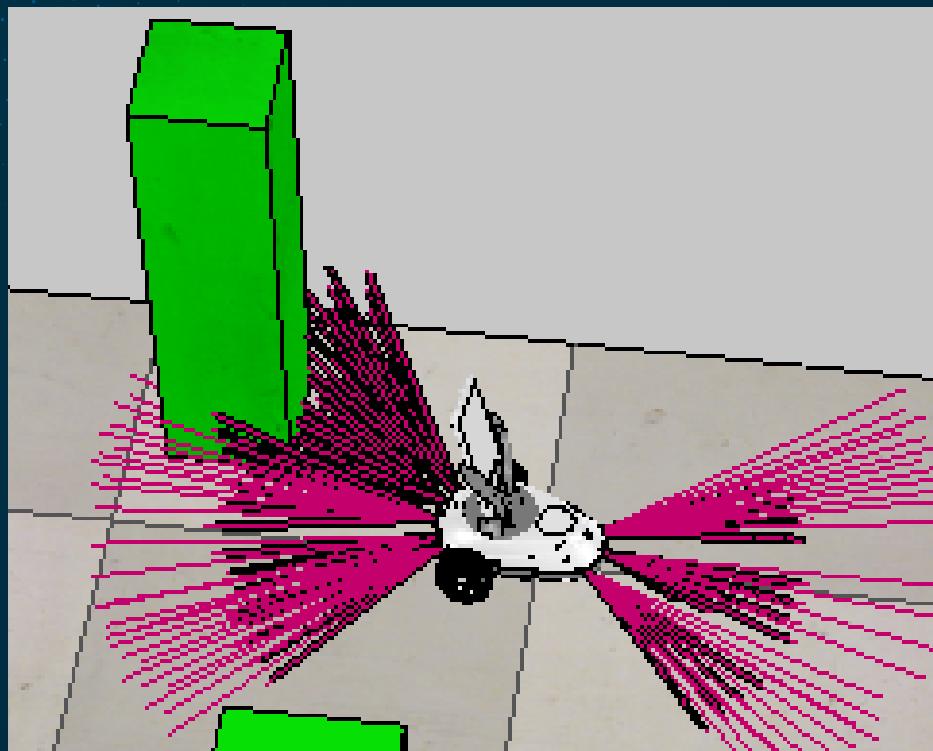


EXPERIMENTAL SETUP

- Three actions ('forward', 'left', and 'right') were used determined by sensor-based state readings.
- The '*Epsilon-Greedy Algorithm*' was used for exploration and exploitation.
- 2 sensors were used: Front Central IR sensor & Mobile Camera.
- 4 sensor-bins were used, and different bin-thresholds for simulation and hardware ([4, 7, 25] [-1, 20, 60]).
- 4 greenness-bins were used, and different bin-thresholds for simulation and hardware ([1, 20, 45] [1, 10, 40]).
- 4 green-direction-bins were used, which takes the max amount of pixels in a direction ([0, 1, 2, 3])
- 50 episodes were used for training, and up to 150 steps were used per episode.
- An *epsilon* of 0.3 was used, allowing for 30% random action selection.
- *Training:*
 - Conducted in simulation.
- *Evaluation:*
 - Total number of collected items per episode to measure the effectiveness of navigational decisions.



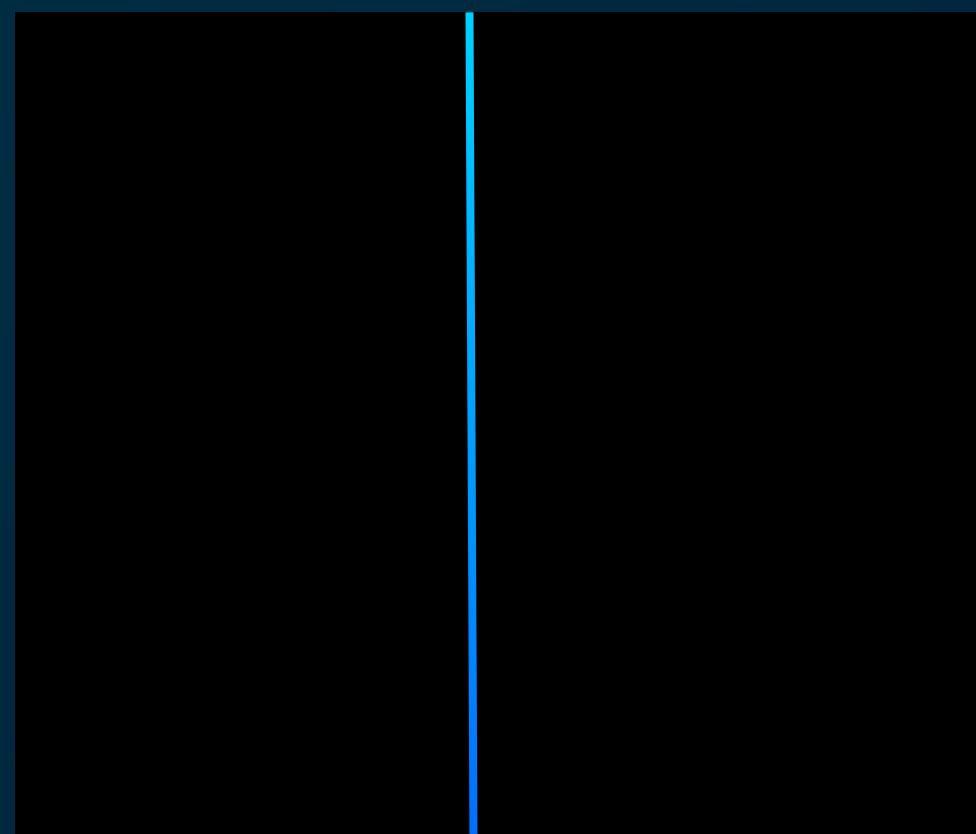
EXPERIMENTAL SETUP



IR value:
9

Green pixels:
[0, 16509, 27233]

25% green in center



State:
[2, 2, 3]

Q-values for this state:
[24.5..., 286.7..., 63.9...]

Go forward!

EXPERIMENTAL SETUP

- **Challenges Encountered:**

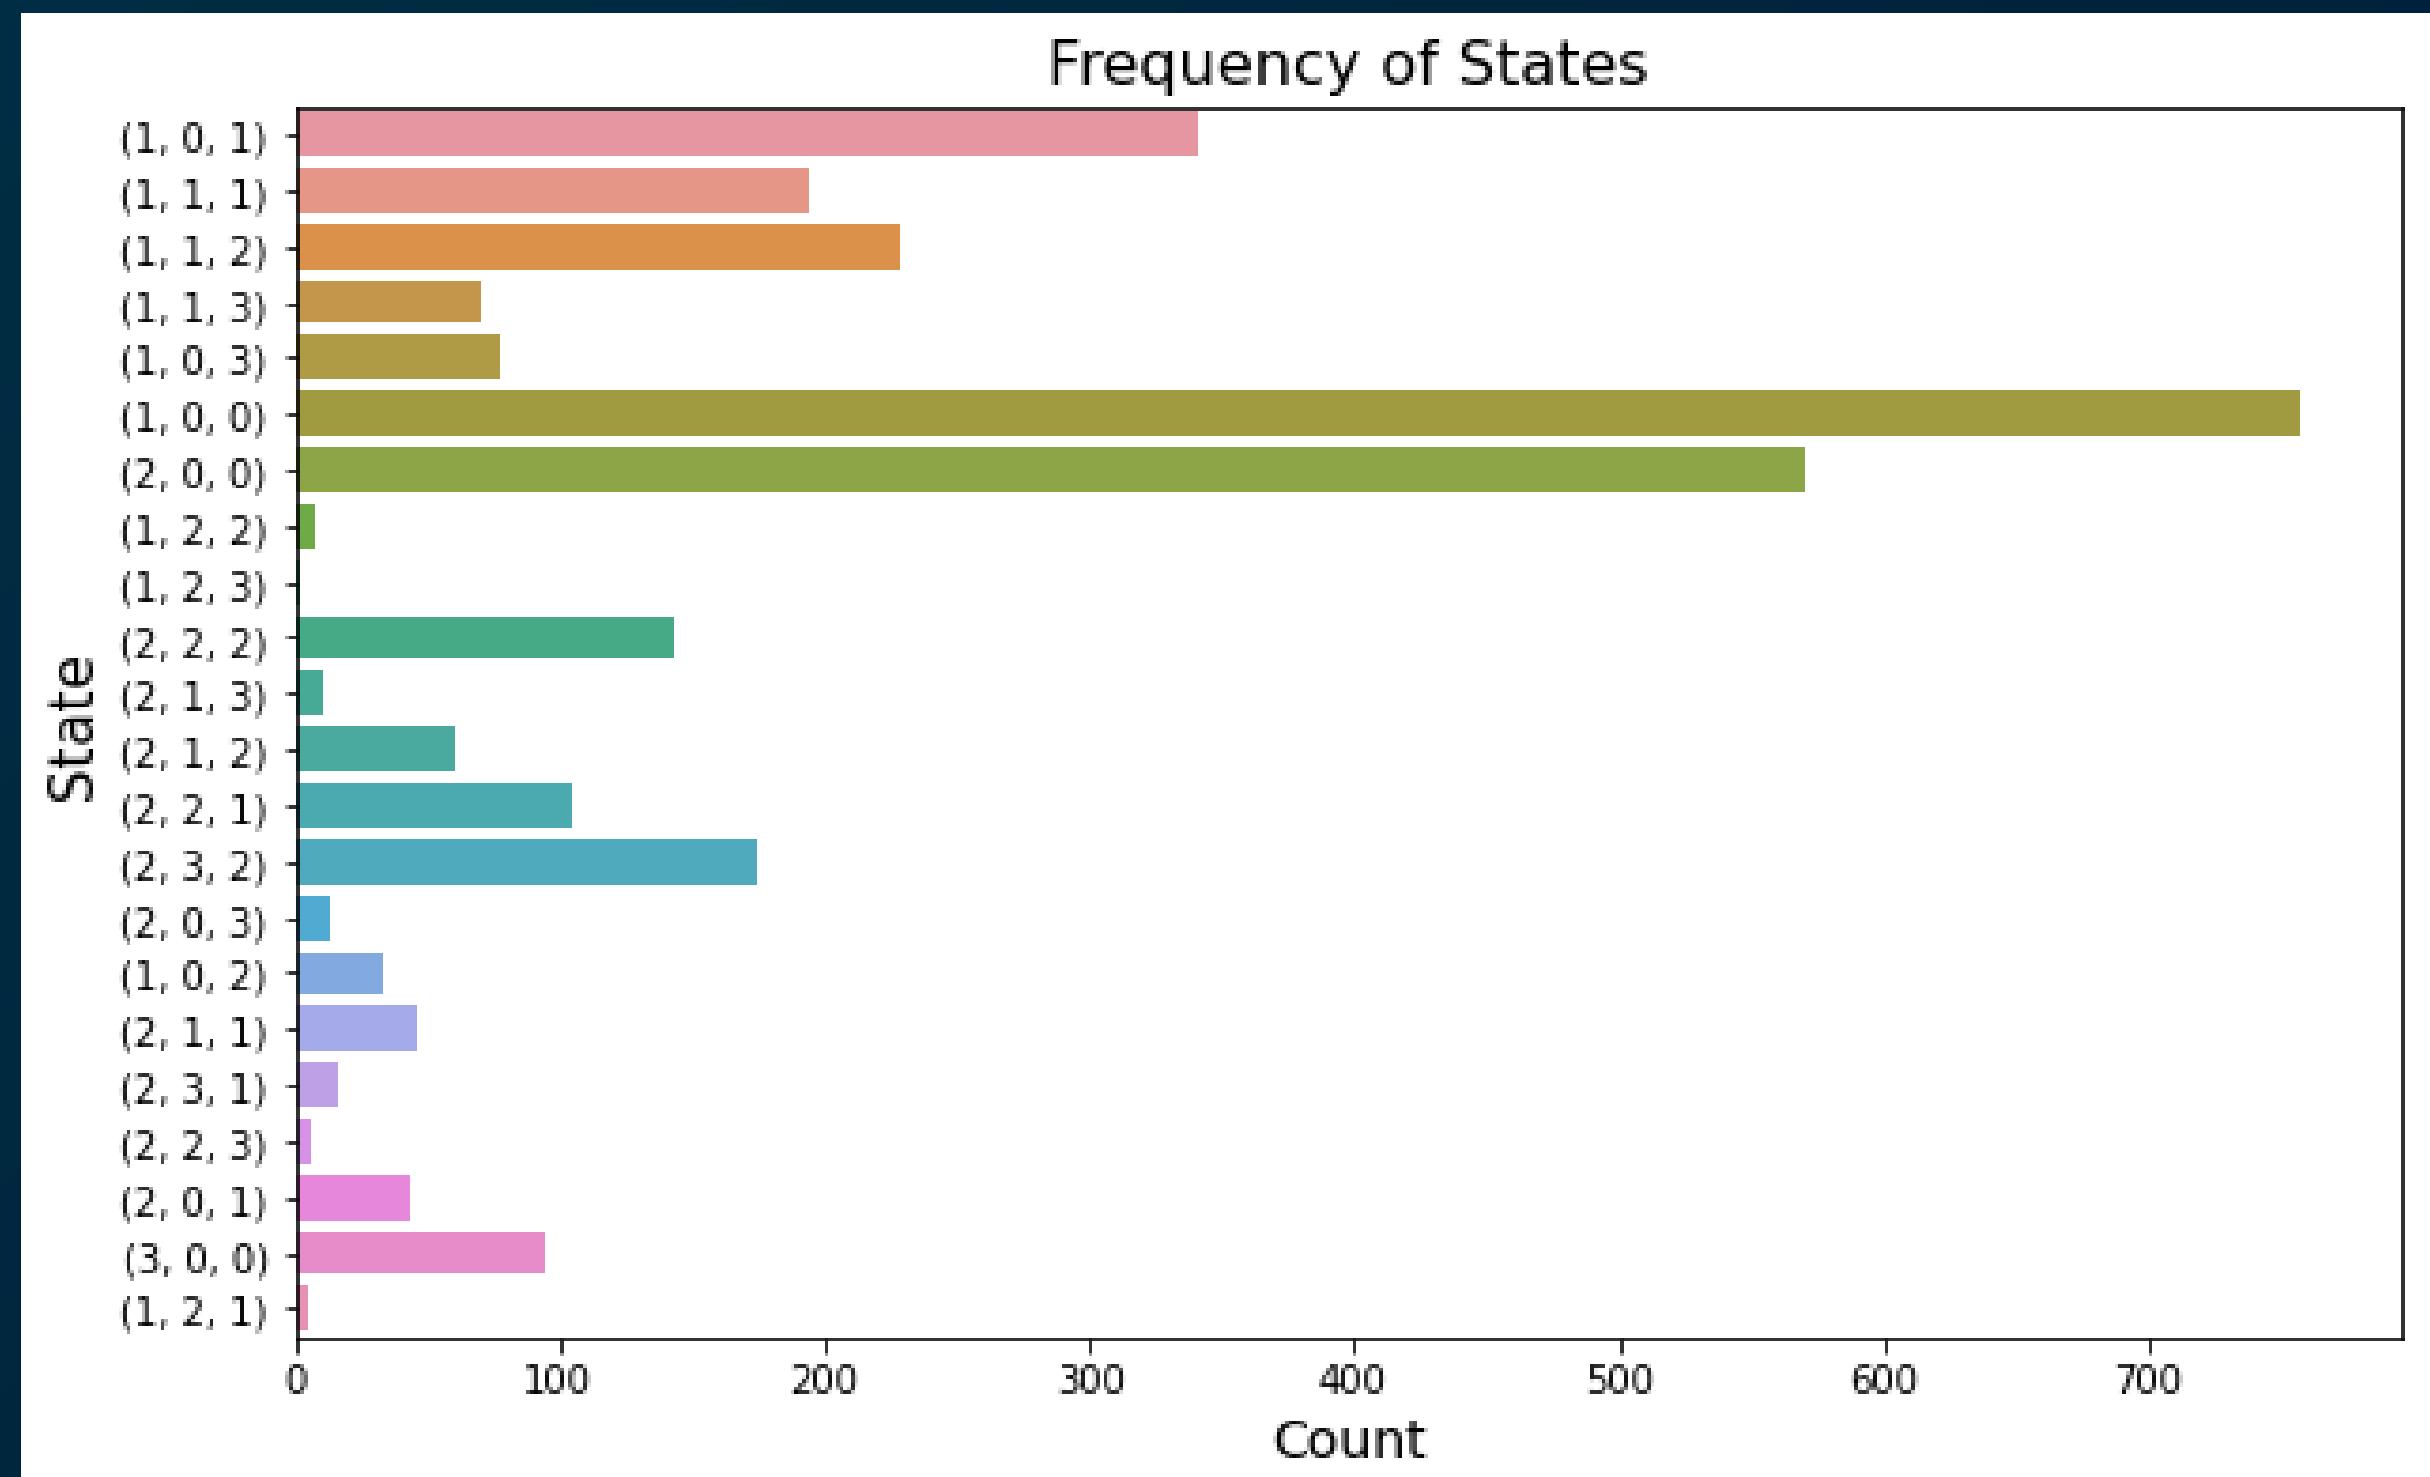
- Determine an appropriate reward function.
- Using only 1 sensor → robot gets stuck in wall often.
- Reality gap:
 - Finding a green pixel filter is much harder in reality.
 - Robot leans slightly to the left when moving forward.

- **Solution Implemented:**

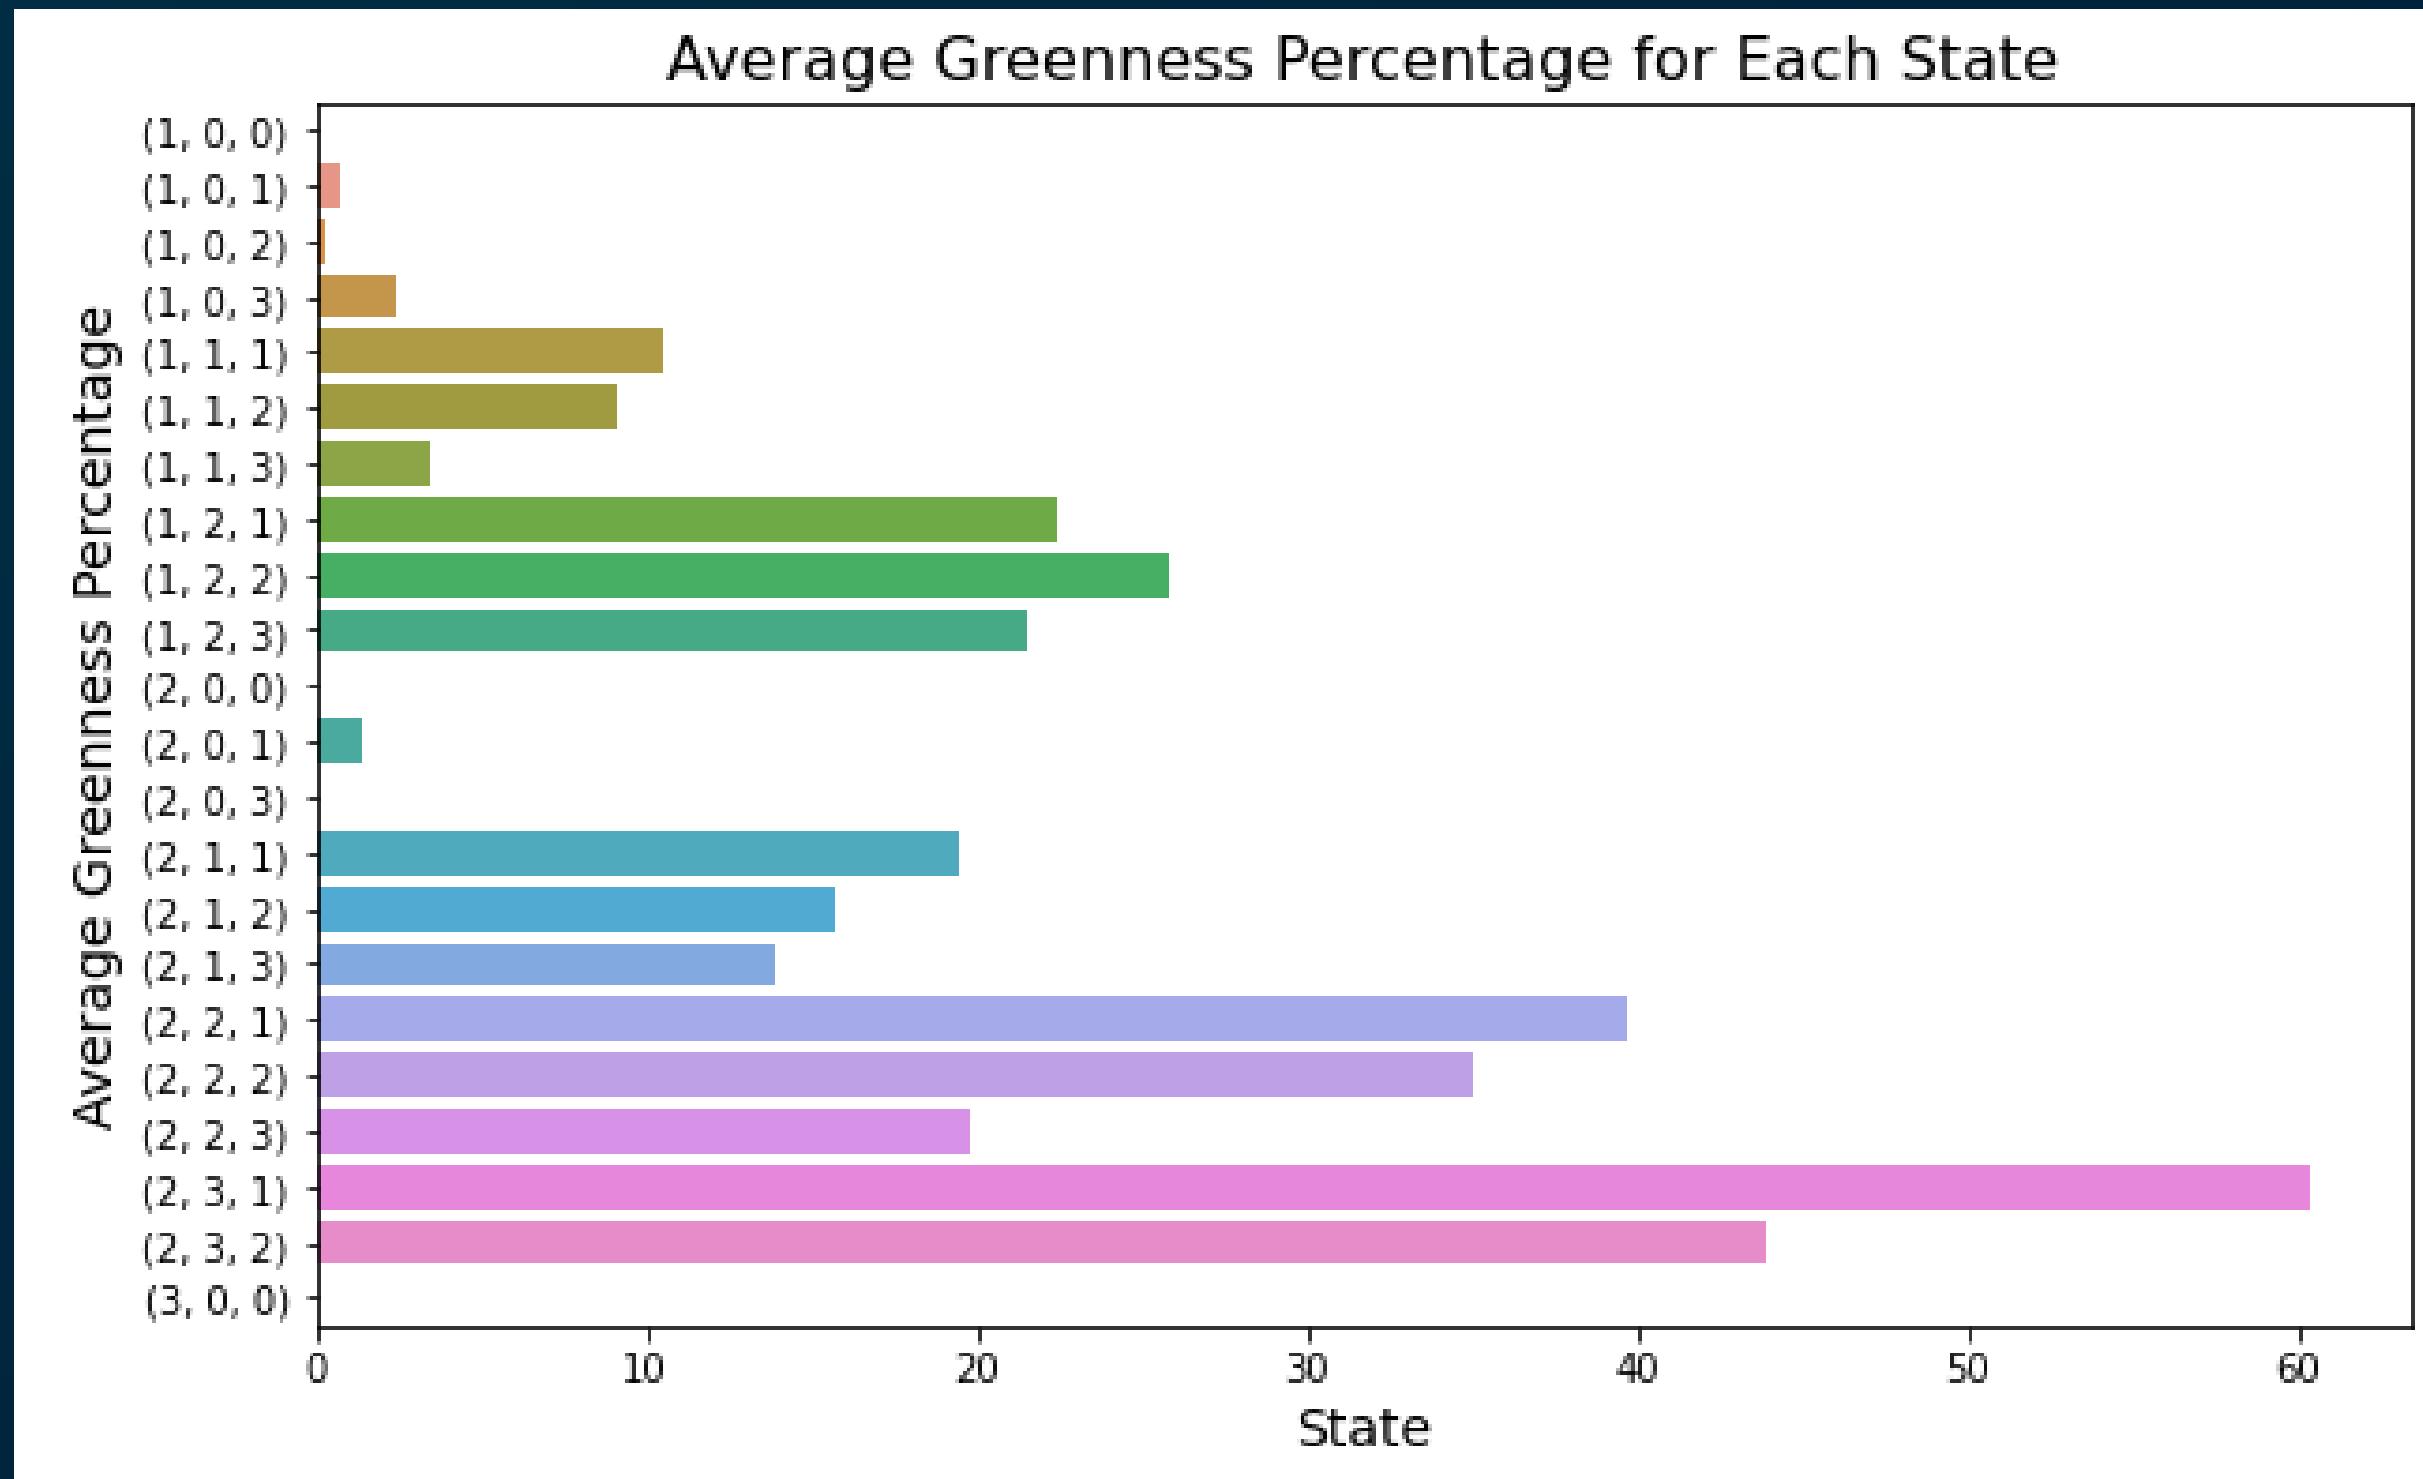
- Trial and error.
- Alter reward function to motivate robot to steer away when stuck.
- Experiment with RGB and camera to see what the robot actually picks up.



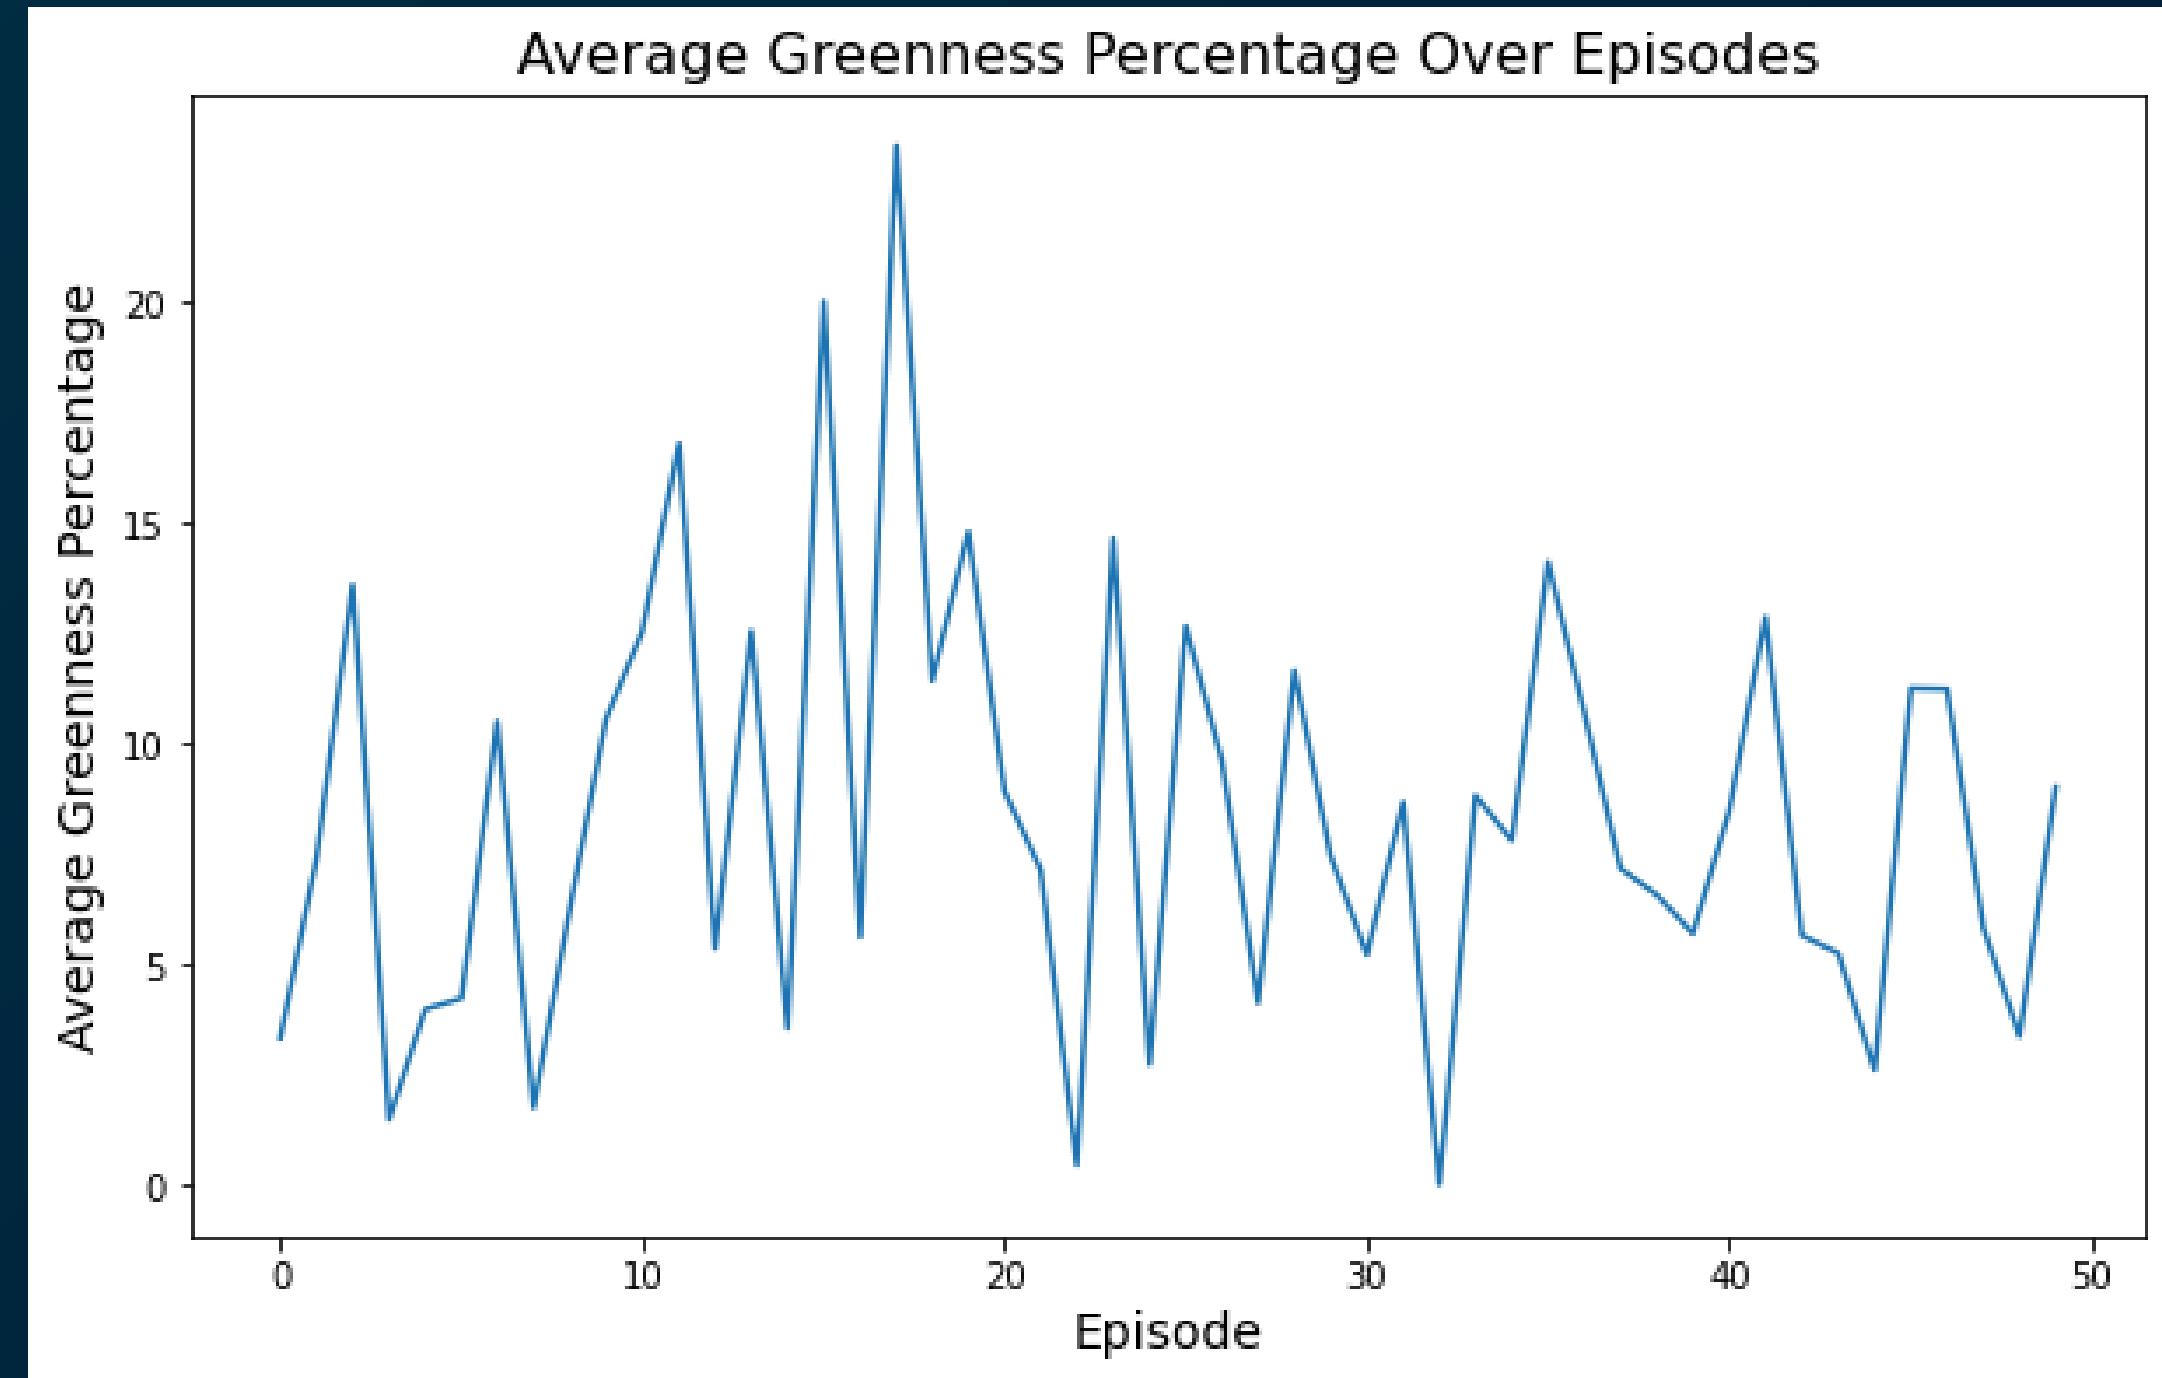
RESULTS: WHAT STATES ?



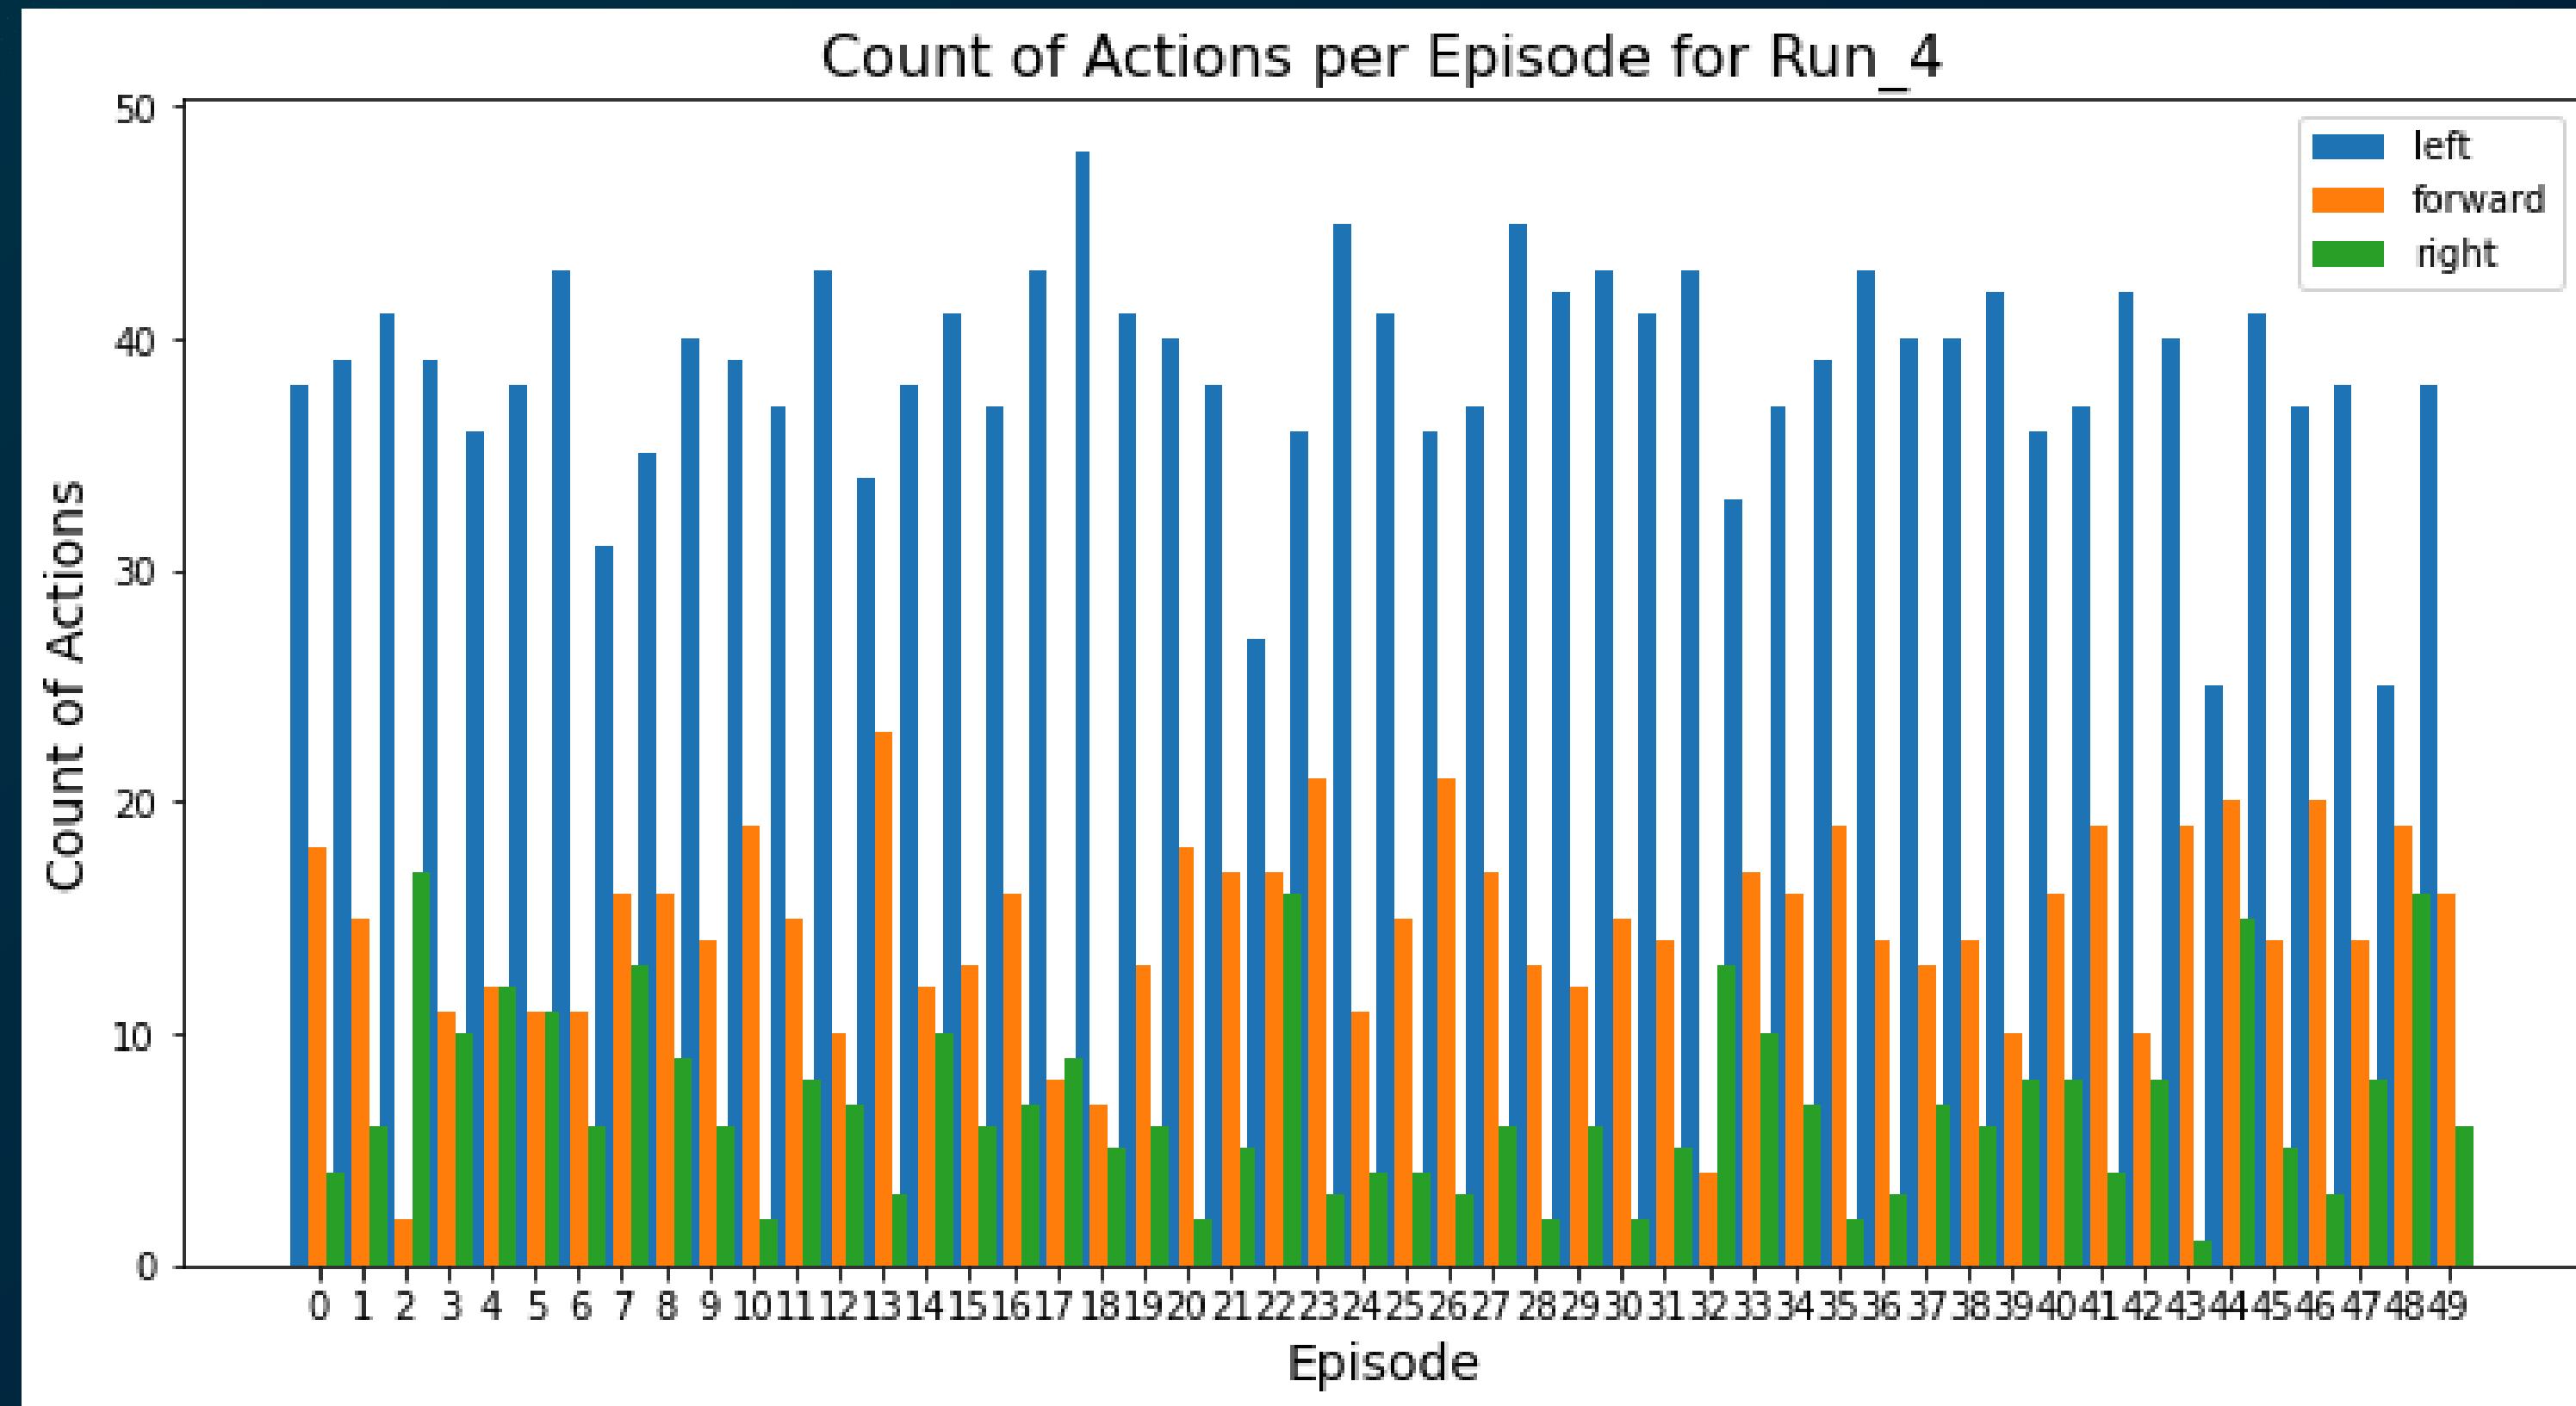
RESULTS: DO WE DETECT GREENNESS?



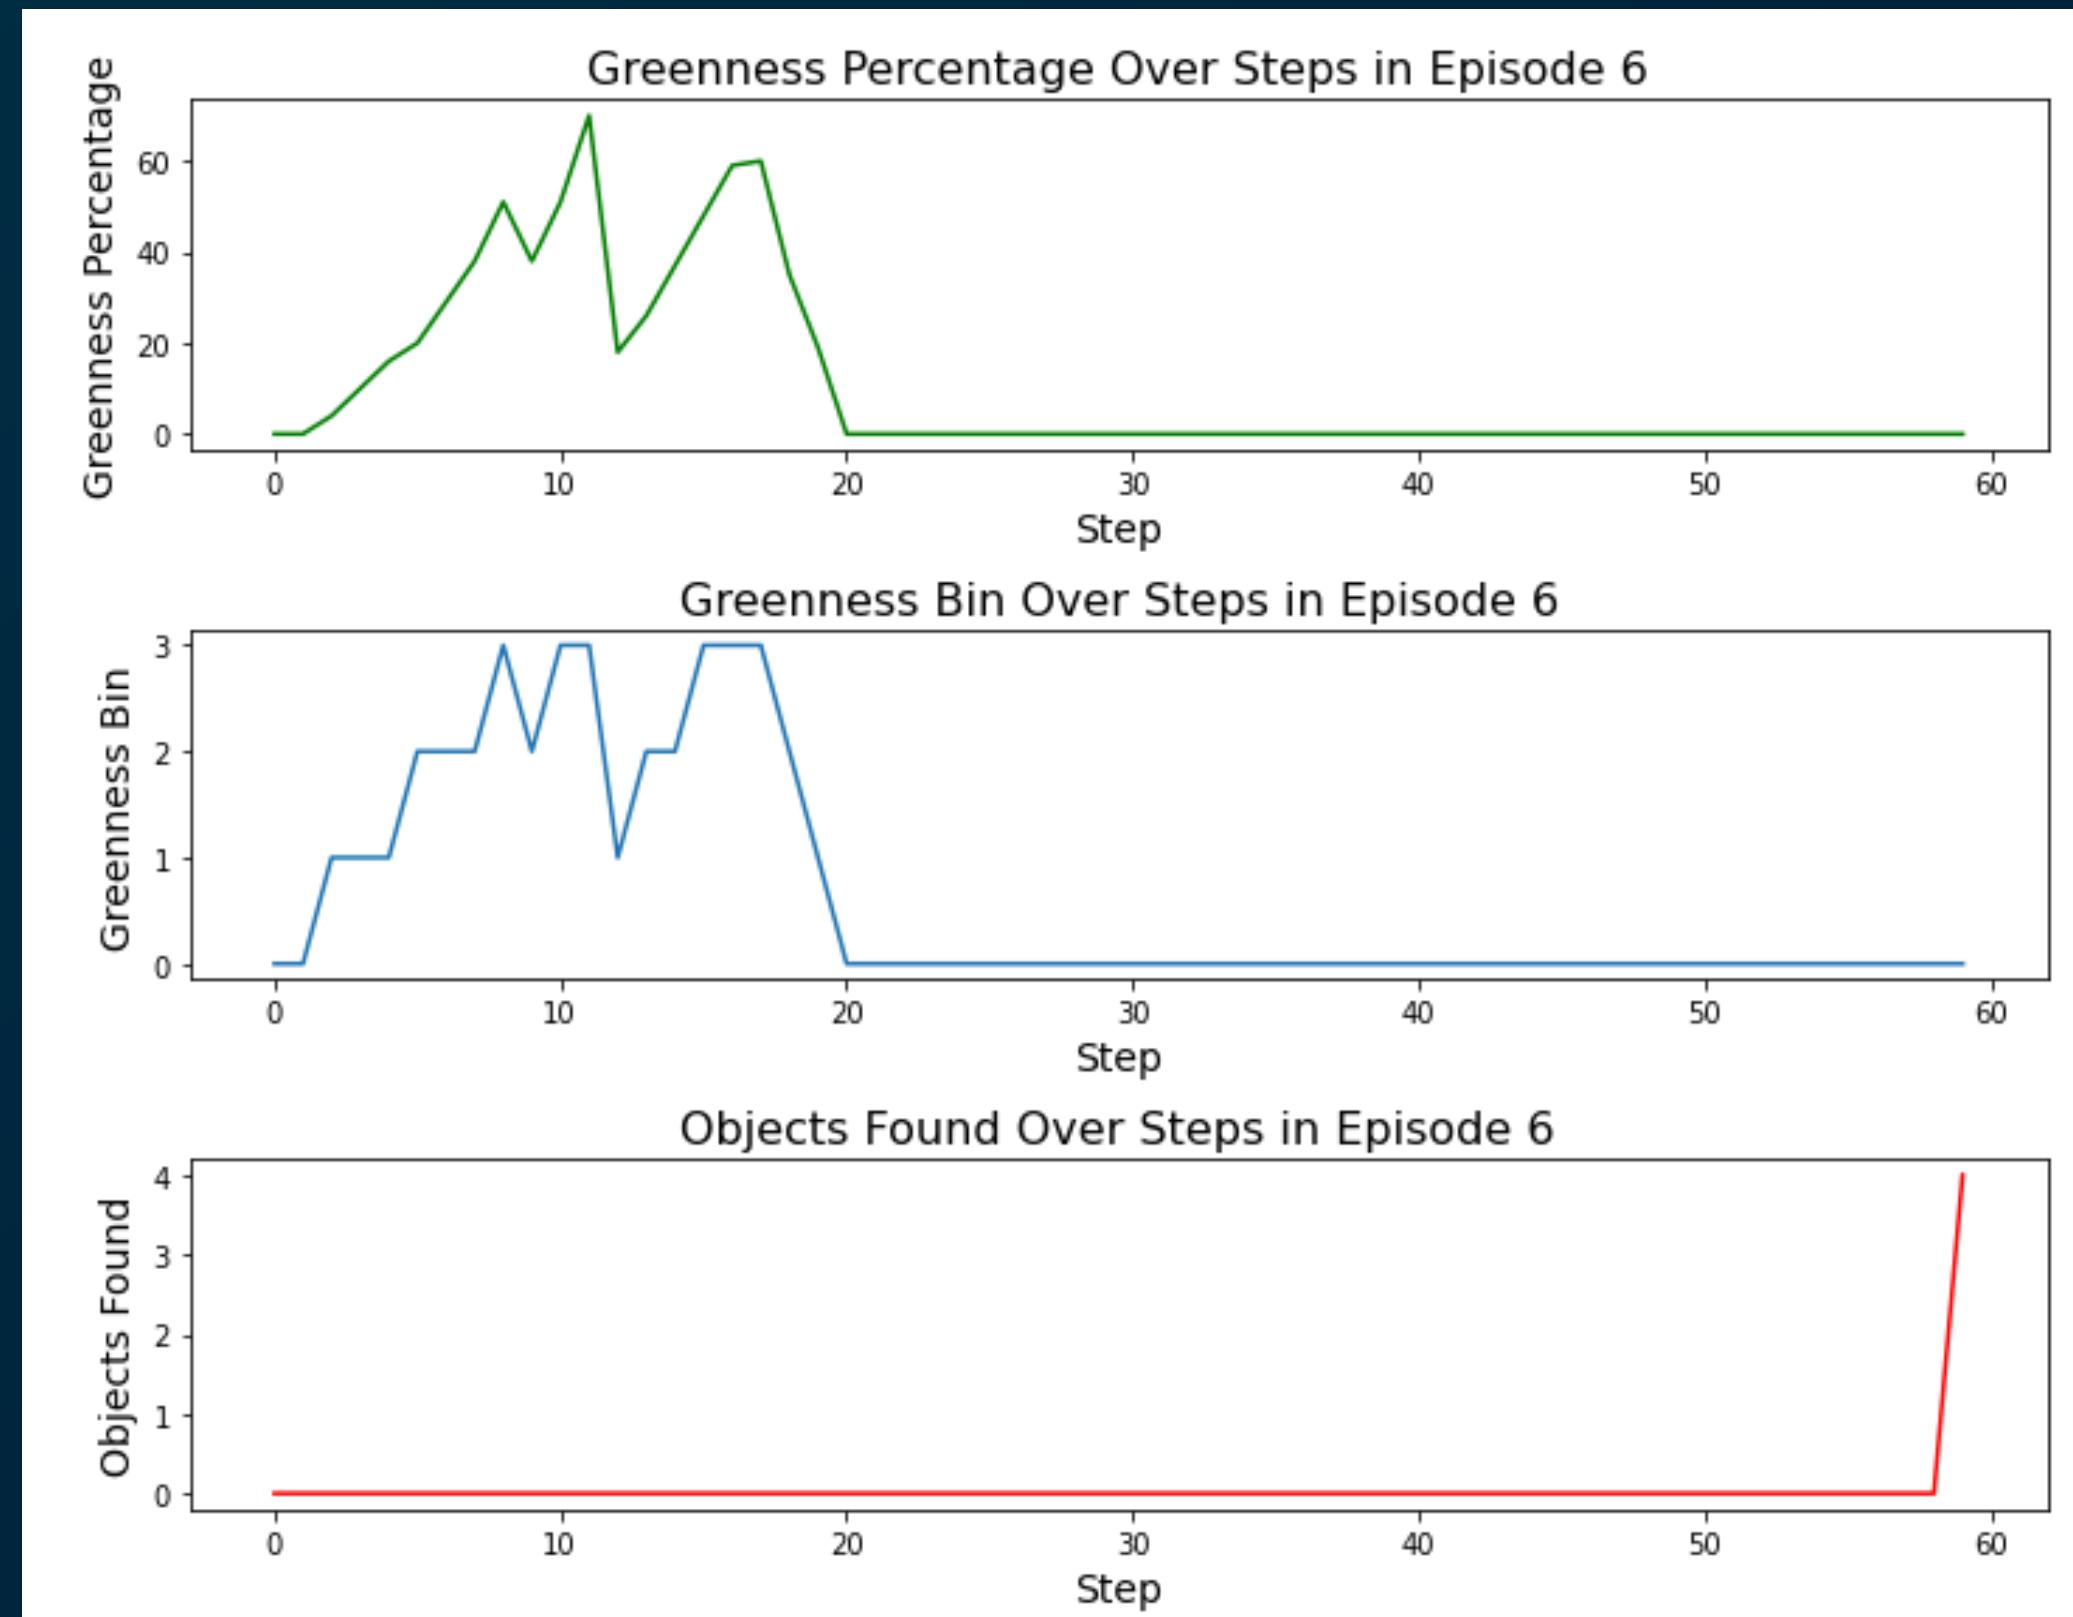
RESULTS: HOW FAST DO WE LEARN?



RESULTS: HOW FAST DO WE LEARN?



RESULTS: HOW FAST DO WE LEARN?



RESULTS: DO WE UNDERSTAND Q-TABLE?

| State | Q-values (left, forward, right) |
|-----------|---|
| (1, 0, 0) | [89.12125342284683, 89.46763971700194, 79.54782758554464] |
| (1, 0, 1) | [189.9621171481047, 132.28667653201498, 73.31675587217163] |
| (1, 0, 2) | [67.54000128837055, 132.06308072521648, 36.38654302973711] |
| (1, 0, 3) | [86.86689297179747, 107.29163940066005, 186.0146593288292] |
| (1, 1, 0) | [0.0, 0.0, 0.0] |
| (1, 1, 1) | [257.5011812105793, 136.71256836524378, 128.81471541970532] |
| (1, 1, 2) | [106.3165210430226, 222.79536969759133, 143.14049038045465] |
| (1, 1, 3) | [37.09893972321737, 70.37274988961495, 180.73849954452137] |
| (1, 2, 0) | [0.0, 0.0, 0.0] |
| (1, 2, 1) | [273.10232825217696, 60.34970107310593, 40.01060276725058] |
| (1, 2, 2) | [54.43743412250306, 279.6672515946146, 95.54361230177116] |
| (1, 2, 3) | [10.407063501407588, 67.2478242616412, 0.0] |
| (1, 3, 0) | [0.0, 0.0, 0.0] |
| (1, 3, 1) | [51.53408731337623, 0.0, 0.0] |
| (1, 3, 2) | [0.0, 0.0, 0.0] |
| (1, 3, 3) | [0.0, 0.0, 0.0] |
| (2, 0, 0) | [61.63453784627899, -14.8447986810351, -2.210094243525016] |
| (2, 0, 1) | [156.33720459509175, 7.324134771738009, -8.290241499933833] |
| (2, 0, 2) | [0.0, 0.0, 0.0] |
| (2, 0, 3) | [35.01709326465945, 5.846646319521613, 284.9003000029957] |
| (2, 1, 0) | [0.0, 0.0, 0.0] |
| (2, 1, 1) | [192.79617837520652, 21.67534276772527, 0.8108334836484261] |
| (2, 1, 2) | [3.2754461715166485, 49.91624708536937, 0.0] |
| (2, 1, 3) | [25.487831672538796, 71.53630101240147, 260.6709883301907] |

RESULTS: DO WE UNDERSTAND Q-TABLE?

| | |
|-----------|--|
| (2, 2, 0) | [0.0, 0.0, 0.0] |
| (2, 2, 1) | [283.16683568799914, 87.91964295990556, 54.58849035032539] |
| (2, 2, 2) | [70.44471036002687, 320.8118211950832, 136.9528844160416] |
| (2, 2, 3) | [26.966628435003297, 294.33854188333146, 71.53737097227676] |
| (2, 3, 0) | [0.0, 0.0, 0.0] |
| (2, 3, 1) | [140.43315005753897, 209.50662239682515, 72.36459722085758] |
| (2, 3, 2) | [182.0355949898773, 368.36090117464323, 142.2643422189268] |
| (2, 3, 3) | [42.52074569317259, 285.7283503754035, 20.48913294569919] |
| (3, 0, 0) | [5.4379576766746505, -79.87224433089195, -73.57451164753292] |
| (3, 0, 1) | [128.80397211011314, 0.0, -9.615251974570143] |
| (3, 0, 2) | [0.0, 0.0, 0.0] |
| (3, 0, 3) | [4.994406043921125, 4.133053275373953, 220.74413125015505] |
| (3, 1, 0) | [0.0, 0.0, 0.0] |
| (3, 1, 1) | [139.11203364554535, -25.18917009933058, 0.0] |
| (3, 1, 2) | [0.0, 0.0, 0.0] |
| (3, 1, 3) | [25.052324408838288, -3.0225648174040014, 260.8269640804169] |
| (3, 2, 0) | [0.0, 0.0, 0.0] |
| (3, 2, 1) | [208.73231477947195, -20.179334204373788, -3.3729855567160003] |
| (3, 2, 2) | [0.0, 0.0, 0.0] |
| (3, 2, 3) | [68.75783694740456, 96.04778376421241, 324.2768778326491] |
| (3, 3, 0) | [0.0, 0.0, 0.0] |
| (3, 3, 1) | [390.18024984613885, 136.20158696582, 186.38894591677968] |
| (3, 3, 2) | [302.10761222614343, 212.25947739248357, 150.01411672018912] |
| (3, 3, 3) | [257.8468905958769, 198.72657463694443, 384.28452063352194] |

CONCLUSIONS

- There is not really a difference between the number of bins chosen.
- Hardware works with selecting right Bin thresholds.
 - Could be done automatically by first calibrating sensors, then mapping to bin.
 - Too much time consuming.
- Future work:
 - Try to implement a Deep RL algorithm.



THANK YOU

QUESTIONS ???



CONTRIBUTION SHEET

- **Stergios:**

- Created the structure and content for the presentation slides.
- Experimented with the code for training the robot.
- Experimented with different algorithms (CNNs).

- **Thijs:**

- Improved on code.
- Experimented with IR values.
- Helped design and experiment with reward function.
- Experimented with hardware.

- **Quirinus:**

- Improved on code from T1: building new state, introducing separate functions.
- Created code for storing data while running with a Training_Results class.
- Created plots from saved data.



BUDDY-GROUP SUMMARY

- **Stergios:**

- Better understanding of my team's capabilities.
- Diverse perspectives improved my ability to see multiple sides of a situation before making decisions.
- Learned more about the functionalities of the robot, including camera operation and object detection.

- **Thijs:**

- Diverse Perspectives.
- Learned practical techniques for minimizing the reality gap in simulations.
- Learned how to use camera in RL algorithms.

- **Quirinus:**

- Got better understanding of using Git.
- Learned how to implement Training_Results class.
- Learned to write better simple functions.
- We started working together better with a plan and some task division.

